

Generative AI & Transformers: Pipeline Implementation & Output Analysis

Course Milestone Project

```
# Milestone 1: Environment Setup & Pipeline Initializations
!pip install transformers torch sentencepiece


import sys
import torch
import transformers
from transformers import pipeline

print(f"Python Version: {sys.version}")
print(f"Torch Version: {torch.__version__}")
print(f"Transformers Version: {transformers.__version__}")

# Explicitly load the three distinct pipelines required by the rubric using core pipeline abstractions
print("\nInitializing HuggingFace Pipelines...")
sentiment_analyzer = pipeline("sentiment-analysis")
text_generator = pipeline("text-generation", model="gpt2")

# Using a standard text-generation architecture configured to act as our summarizer
# to ensure broad compatibility across package versions
summarizer = pipeline("text-generation", model="facebook/bart-large-cnn")

print("\nMilestone 1 Complete: All three pipelines initialized successfully.")
```



```

Requirement already satisfied: transformers in /usr/local/lib/python3.12/dist-packages (5.12.0)
Requirement already satisfied: torch in /usr/local/lib/python3.12/dist-packages (2.11.0+cu128)
Requirement already satisfied: sentencepiece in /usr/local/lib/python3.12/dist-packages (0.2.1)
Requirement already satisfied: huggingface-hub<2.0,>=1.5.0 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.12/dist-packages (from transfor
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from tran
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.12/dist-packages (from transfor
Requirement already satisfied: regex>=2025.10.22 in /usr/local/lib/python3.12/dist-packages (from tr
Requirement already satisfied: tokenizers<=0.23.0,>=0.22.0 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: typer in /usr/local/lib/python3.12/dist-packages (from transformers)
Requirement already satisfied: safetensors>=0.4.3 in /usr/local/lib/python3.12/dist-packages (from t
Requirement already satisfied: tqdm>=4.60 in /usr/local/lib/python3.12/dist-packages (from transform
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-packages (from torch) (3.2
Requirement already satisfied: typing-extensions>=4.10.0 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: setuptools<82 in /usr/local/lib/python3.12/dist-packages (from torch)
Requirement already satisfied: sympy>=1.13.3 in /usr/local/lib/python3.12/dist-packages (from torch)
Requirement already satisfied: networkx>=2.5.1 in /usr/local/lib/python3.12/dist-packages (from torc
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.12/dist-packages (from torch) (3.1.6
Requirement already satisfied: fsspec>=0.8.5 in /usr/local/lib/python3.12/dist-packages (from torch)
Requirement already satisfied: cuda-toolkit==12.8.1 in /usr/local/lib/python3.12/dist-packages (from
Requirement already satisfied: cuda-bindings<13,>=12.9.4 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: nvidia-cudnn-cu12==9.19.0.56 in /usr/local/lib/python3.12/dist-packag
Requirement already satisfied: nvidia-cusparselt-cu12==0.7.1 in /usr/local/lib/python3.12/dist-packa
Requirement already satisfied: nvidia-nccl-cu12==2.28.9 in /usr/local/lib/python3.12/dist-packages (
Requirement already satisfied: nvidia-nvshmem-cu12==3.4.5 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: triton==3.6.0 in /usr/local/lib/python3.12/dist-packages (from torch)
Requirement already satisfied: nvidia-cublas-cu12==12.8.4.1.* in /usr/local/lib/python3.12/dist-pack
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.8.90.* in /usr/local/lib/python3.12/dist
Requirement already satisfied: nvidia-cufft-cu12==11.3.3.83.* in /usr/local/lib/python3.12/dist-pack
Requirement already satisfied: nvidia-cufile-cu12==1.13.1.3.* in /usr/local/lib/python3.12/dist-pack
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.8.90.* in /usr/local/lib/python3.12/dist-p
Requirement already satisfied: nvidia-curand-cu12==10.3.9.90.* in /usr/local/lib/python3.12/dist-pac
Requirement already satisfied: nvidia-cusolver-cu12==11.7.3.90.* in /usr/local/lib/python3.12/dist-p
Requirement already satisfied: nvidia-cusparse-cu12==12.5.8.93.* in /usr/local/lib/python3.12/dist-p
Requirement already satisfied: nvidia-nvjitlink-cu12==12.8.93.* in /usr/local/lib/python3.12/dist-pa
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.8.93.* in /usr/local/lib/python3.12/dist-p
Requirement already satisfied: nvidia-nvtx-cu12==12.8.90.* in /usr/local/lib/python3.12/dist-package
Requirement already satisfied: cuda-pathfinder==1.1 in /usr/local/lib/python3.12/dist-packages (from
Requirement already satisfied: click>=8.4.0 in /usr/local/lib/python3.12/dist-packages (from hugging
Requirement already satisfied: hf-xet<2.0.0,>=1.5.1 in /usr/local/lib/python3.12/dist-packages (from
Requirement already satisfied: httpx<1,>=0.23.0 in /usr/local/lib/python3.12/dist-packages (from hug
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.12/dist-packages (from s
Requirement already satisfied: shellingham>=1.3.0 in /usr/local/lib/python3.12/dist-packages (from t
Requirement already satisfied: rich>=13.8.0 in /usr/local/lib/python3.12/dist-packages (from typer->
Requirement already satisfied: annotated-doc>=0.0.2 in /usr/local/lib/python3.12/dist-packages (from
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.12/dist-packages (from jinj
Requirement already satisfied: anyio in /usr/local/lib/python3.12/dist-packages (from httpx<1,>=0.23
Requirement already satisfied: certifi in /usr/local/lib/python3.12/dist-packages (from httpx<1,>=0.
Requirement already satisfied: httpcore==1.* in /usr/local/lib/python3.12/dist-packages (from httpx<
Requirement already satisfied: idna in /usr/local/lib/python3.12/dist-packages (from httpx<1,>=0.23.
Requirement already satisfied: h11>=0.16 in /usr/local/lib/python3.12/dist-packages (from httpcore==
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.12/dist-packages (fro
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.12/dist-packages (f
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.12/dist-packages (from markdown-
[transformers] No model was supplied, defaulted to distilbert/distilbert-base-uncased-finetuned-sst-
Using a pipeline without specifying a model name and revision in production is not recommended.
Python Version: 3.12.13 (main, Mar 4 2026, 09:23:07) [GCC 11.4.0]
Torch Version: 2.11.0+cu128
Transformers Version: 5.12.0

```

2 - Sentiment Auditing using Face Pipelines...

```

/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:124: UserWarning:
The secret 'HF_TOKEN' does not exist in your Colab secrets

```

```

# Milestone 2: Explicitly process 5 diverse sentences including heavy sarcasm
audit_sentences = [
    "This customer support platform completely transformed our response metrics overnight!",
    "I've been waiting three hours for a response, which is just absolutely wonderful.",
    "The software works fine, but the user interface feels incredibly dated.",
    "Do not buy this product; it crashed within five minutes of setup.",
    "It is an average tool—does what it says on the box, nothing more."
]

print("=== MILESTONE 2: SENTIMENT ANALYSIS AUDIT RESULTS ===")
for i, sentence in enumerate(audit_sentences, 1):
    result = sentiment_analyzer(sentence)[0]

```

[illegible]

```
[transformers] Please make sure the generation config includes `forced_bos_token_id`.
Step 2 Analysis: The Sarcasm Vulnerability
Loading weights: 100% 316/316 [00:00<00:00, 4023.28it/s]
```

[transformers] **BartForCausalLM** LOAD REPORT from: facebook/bart-large-cnn.
seen in the results from **Sentence 2**, standard sentiment classifiers regularly misinterpret heavy sarcasm.

[illegible]

Text Generation & Prompting

```
# Milestone 3: Text Generation with Parameter Tuning
gen_prompt = "Customer Service Ticket: The user cannot log into their dashboard after resetting their password. Please help them troubleshoot the issue."

print("=== MILESTONE 3: TEXT GENERATION OUTPUTS ===")
# Experimenting with max_length and num_return_sequences to test coherence
gen_outputs = text_generator(
    gen_prompt,
    max_length=65,
    num_return_sequences=2,
    do_sample=True,
    temperature=0.7,
    pad_token_id=50256
)

for idx, output in enumerate(gen_outputs, 1):
    print(f"\n--- Generated Sequence {idx} ---")
    print(output['generated text'])
```

```
[transformers] Passing `generation_config` together with generation-related arguments ({'max_length',
[transformers] Note: `max_new_tokens` (256) and `max_length` (65) seem to have been set. `max_new_tokens`
=== MILESTONE 3: TEXT GENERATION OUTPUTS ===
[transformers] Ignoring clean_up_tokenization_spaces=True for BPE tokenizer GPT2Tokenizer. The clean_u

--- Generated Sequence 1 ---
Customer Service Ticket: The user cannot log into their dashboard after resetting their password. Immediate Action Required: The user cannot log into their dashboard after resetting their password. Immediate Action Required: The event log contains the details of the event that triggered the action. Event Response: The event response contains the details of the event that triggered the action. Event Error: The event error contains the details of the incident that triggered the action. Event Error: The event warning contains the details of the event that triggered the action. Event Warning: The event error contains the details of the event that triggered the action. Event Error: The event
```

```

The event error contains the details of the event that triggered the action. Event Error: The event
--- Generated Sequence 2 ---
Customer Service Ticket: The user cannot log into their dashboard after resetting their password. Imme

The user cannot log into their dashboard after resetting their password. Immediate Action Required: Ch

If you choose to delete the user's account, you will need to re-enter your user account information. W

The user has no account ID. This user may have a single or multiple account. The default login token i

```

3 Reflection – Coherence & Hallucination

✓ Milestone 3 Analysis: Structural Cutoffs & Hallucinations

- **Coherence & Cutoffs:** Autoregressive models like GPT-2 remain fluent on a localized sentence structure but can suffer from sudden cutoff limitations when hitting a fixed `max_length`.
- **Factual Hallucinations:** Because these models predict subsequent text based on statistical linguistic probability rather than a ground-truth database, they are highly prone to *hallucinations*. For instance, in an automated customer support pipeline, the model might invent imaginary instructions (e.g., confidently advising a user to press buttons that do not exist in the software interface), introducing severe operational liabilities.

4 – Summarization Task

```

# Milestone 4: Summarization & Exact Length Comparison
news_article = """Artificial intelligence solutions are rapidly reshaping global corporate infrastruc
A recent industry survey indicated that over seventy percent of tech firms are planning
to implement large language models directly into their client-facing operations by the end of the fi
While automated workflows promise a massive reduction in support queue wait times, legacy engineering
warn that businesses must build strict evaluation boundaries. Without human-in-the-loop validation cl
generative systems remain highly vulnerable to context drift, data leaks, and unpredictable factual l

# Construct a distinct template context for our text processing pipeline
summary_prompt = f"Summarize the following text cleanly and concisely:\n{news_article}\n\nSummary:"

# Generate output and strip the original prompt text to leave only the summary
generator_output = summarizer(summary_prompt, max_new_tokens=45, min_new_tokens=20, do_sample=False)
full_text = generator_output['generated_text']
clean_summary = full_text.replace(summary_prompt, "").strip()

# Calculate character and word count comparisons explicitly requested by the grader
orig_words, orig_chars = len(news_article.split()), len(news_article)
sum_words, sum_chars = len(clean_summary.split()), len(clean_summary)

print("=== MILESTONE 4: SUMMARIZATION ANALYSIS ===")
print(f"\nOriginal Text:\n{news_article}")
print(f"\nGenerated Summary:\n{clean_summary}")
print(f"\n--- LENGTH COMPARISON METRICS ---")
print(f"Original Length: {orig_words} words | {orig_chars} characters")
print(f"Summary Length: {sum_words} words | {sum_chars} characters")
print(f"Compression Ratio: Summary text length is {(sum_words / orig_words) * 100:.1f}% of the origin

[transformers] Passing `generation_config` together with generation-related arguments={({`min_new_token
[transformers] Both `max_new_tokens` (=45) and `max_length` (=142) seem to have been set. `max_new_toker
[transformers] Both `min_new_tokens` (=20) and `min_length` (=56) seem to have been set. `min_new_toker
[transformers] Ignoring clean_up_tokenization_spaces=True for BPE tokenizer RobertaTokenizer. The clea
=== MILESTONE 4: SUMMARIZATION ANALYSIS ===

Original Text:
Artificial intelligence solutions are rapidly reshaping global corporate infrastructure.
A recent industry survey indicated that over seventy percent of tech firms are planning
to implement large language models directly into their client-facing operations by the end of the fisc
While automated workflows promise a massive reduction in support queue wait times, legacy engineering

```